



SALONI SINGH

CSE(AI)-C

202401100300211

Iris Flower Classification

INTRODUCTION

The **Iris Flower Classification** is a classic machine learning problem where the goal is to classify iris flowers into three species:

1. **Setosa**
2. **Versicolor**
3. **Virginica**

based on four numerical features:

- **Sepal Length** (cm)
- **Sepal Width** (cm)
- **Petal Length** (cm)
- **Petal Width** (cm)

Steps in Classification

1. Data Loading & Preprocessing

- Read the dataset using **pandas**.
- Check for missing or incorrect values.
- Encode the species column into numerical labels.

2. Exploratory Data Analysis (EDA)

- Visualize feature relationships using scatter plots and pair plots.
- Check class distribution and correlation between features.

3. Train-Test Split

- Split the dataset into **80% training** and **20% testing** data.
- Use stratified sampling to maintain class distribution.

4. Model Selection & Training

- Train classifiers like **Random Forest, SVM, or KNN**.
- Tune hyperparameters for optimal performance.

5. Evaluation

- Use **accuracy score, confusion matrix, and classification report**.
- Visualize model performance using heatmaps and feature importance graphs.

Why is it Important?

The Iris dataset is widely used for:

- Learning **classification algorithms**.
- Understanding **data preprocessing and visualization**.
- Practicing **model evaluation and optimization**.

METHODOLOGY

The **methodology** for classifying iris flowers involves several key steps, from data collection to model evaluation. Below is a structured approach to solving the problem:

1. Data Collection & Understanding

- The dataset is sourced from **UCI Machine Learning Repository** and contains **150 samples** with **4 features** (sepal & petal length/width) and **1 target variable** (species).
 - The three species to classify are **Setosa, Versicolor, and Virginica**.
-

2. Data Preprocessing

- **Load the dataset** using **pandas**.
- **Check for missing values** and handle any inconsistencies.
- **Encode categorical labels** (species) into numerical values using **LabelEncoder**.

3. Exploratory Data Analysis (EDA)

- Use **descriptive statistics** to summarize the data.
- **Visualize relationships** using `seaborn` pair plots and histograms.
- Check for **class imbalances** or feature correlations.

4. Splitting the Dataset

- Divide the dataset into **training (80%)** and **testing (20%)** using `train_test_split()`.
- Use **stratified sampling** to maintain equal class distribution.

5. Model Selection & Training

- Choose classification models like:
 - **Random Forest** (Ensemble learning)
 - **Support Vector Machine (SVM)** (Margin-based classification)
 - **K-Nearest Neighbors (KNN)** (Distance-based classification)
 - Train the models on the **training dataset** using `fit()`.
-

6. Model Evaluation

- Make predictions on the **test dataset** using `predict()`.
 - Assess model performance using:
 - **Accuracy Score** – Measures overall correctness.
 - **Confusion Matrix** – Evaluates false positives/negatives.
 - **Classification Report** – Analyzes precision, recall, and F1-score.
 - Visualize performance using heatmaps and feature importance plots.
-

7. Model Optimization

- Tune hyperparameters using **GridSearchCV** or **RandomizedSearchCV**.
 - Experiment with different models to find the best-performing classifier.
-

8. Conclusion & Interpretation

- Summarize findings and compare models.
- Discuss the best model for classifying iris species based on accuracy and performance metrics.

CODE OF THE PROGRAM

```
import pandas as pd

import seaborn as sns

import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import LabelEncoder

from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import accuracy_score,
classification_report, confusion_matrix


# Load the dataset

file_path = '/content/iris_data.csv'


df = pd.read_csv(file_path)

print("Original Column Names:", df.columns)
```

```
# Ensure 'species' column exists
```

```
expected_columns = {'sepal_length', 'sepal_width',  
'petal_length', 'petal_width', 'species'}
```

```
if not expected_columns.issubset(df.columns):
```

```
    print("Warning: Expected columns not found. Checking  
possible issues...")
```

```
    df = pd.read_csv(file_path, header=None)
```

```
    df.columns = ['sepal_length', 'sepal_width', 'petal_length',  
'petal_width', 'species']
```

```
    print("Columns renamed to:", df.columns)
```

```
# Display basic dataset info
```

```
print("Dataset Overview:\n", df.head())
```

```
print("\nDataset Info:\n")
```

```
df.info()
```

```
print("\nSummary Statistics:\n", df.describe())
```

```
# Encode the target variable
```

```
label_encoder = LabelEncoder()
```



```
df['species'] = label_encoder.fit_transform(df['species'])
```

```
# Check class distribution before splitting
```

```
print("\nClass Distribution:\n", df['species'].value_counts())
```

```
# Ensure all classes have at least two samples
```

```
class_counts = df['species'].value_counts()
```

```
if class_counts.min() < 2:
```

```
    print("Warning: Some classes have fewer than 2 samples.  
    Removing them...")
```

```
    df = df[df['species'].map(class_counts) >= 2]
```

```
# Split data into features and target
```

```
X = df.drop(columns=['species'])
```

```
y = df['species']
```

```
# Handle small class issue by removing stratify if needed
```

```
try:
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42, stratify=y)
```

```
except ValueError:
```

```
    print("Stratification failed due to low class counts.  
Proceeding without stratify...")
```

```
    X_train, X_test, y_train, y_test = train_test_split(X, y,  
test_size=0.2, random_state=42)
```

```
# Train a RandomForestClassifier with optimized  
parameters
```

```
model = RandomForestClassifier(n_estimators=200,  
max_depth=5, random_state=42)
```

```
model.fit(X_train, y_train)
```

```
# Make predictions
```

```
y_pred = model.predict(X_test)
```

```
# Evaluate the model
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f'\nModel Accuracy: {accuracy:.2f}')
```

```
print('\nClassification Report:\n',  
classification_report(y_test, y_pred))
```

```
# Confusion Matrix
```

```
conf_matrix = confusion_matrix(y_test, y_pred)
```

```
sns.heatmap(conf_matrix, annot=True, fmt='d',  
cmap='Blues', xticklabels=label_encoder.classes_,  
yticklabels=label_encoder.classes_)
```

```
plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
```

```
plt.title('Confusion Matrix')
```

```
plt.show()
```

```
# Visualizing feature importance
```

```
feature_importances =  
pd.Series(model.feature_importances_, index=X.columns)
```

```
plt.figure(figsize=(8, 5))
```

```
feature_importances.nlargest(4).plot(kind='barh',  
color='skyblue')
```

```
plt.xlabel('Feature Importance')
```

```
plt.ylabel('Features')
```

```
plt.title('Feature Importance in Iris Classification')
```

```
plt.show()
```

OUTPUT

```
Original Column Names: Index(['Sepallength', 'SepalWidth', 'PetalLength', 'PetalWidth', 'Species'], dtype='object')
Warning: Expected columns not found. Checking possible issues...
Columns renamed to: Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
                           'species'],
                           dtype='object')
Dataset Overview:
   sepal_length  sepal_width  petal_length \
0  Sepallength  SepalWidth  PetalLength
1  7.303274553127926  2.4750252534128063  2.1760486154633614
2  7.556927655410556  2.987381008530221  1.9215846218548323
3  5.254016373665658  2.0935160238309827  3.672563875820435
4  6.409620271630198  2.2110415805023966  1.812868616156213

   petal_width  species
0  PetalWidth  Species
1  0.6950029791866986  Setosa
2  1.1726148049564784  Versicolor
3  0.5504235628446024  Virginica
4  1.745372347536642  Versicolor
```

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21 entries, 0 to 20
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   sepal_length  21 non-null    object
1   sepal_width   21 non-null    object
2   petal_length  21 non-null    object
3   petal_width   21 non-null    object
4   species       21 non-null    object
dtypes: object(5)
memory usage: 972.0+ bytes

Summary Statistics:
   sepal_length  sepal_width  petal_length  petal_width  species
count         21          21          21          21         21
unique         21          21          21          21          4
top  Sepallength  SepalWidth  PetalLength  PetalWidth  Setosa
freq           1           1           1           1          7
```

```
Class Distribution:
species
0    7
3    7
2    6
1    1
Name: count, dtype: int64
Warning: Some classes have fewer than 2 samples. Removing them...
```

Model Accuracy: 0.25

```
Classification Report:
              precision    recall  f1-score   support

     0       0.00        0.00     0.00         2
     2       0.33        1.00     0.50         1
     3       0.00        0.00     0.00         1

 accuracy          0.25         0.25         0.25         4
 macro avg          0.11         0.33         0.17         4
 weighted avg          0.08         0.25         0.12         4
```



